Zoe Saldana - Senior MLOps Engineer

Contact

Email: zoe.saldana@mlops.io | Phone: (555) 060-5566 | Location: San Jose, CA

Summary

Senior MLOps Engineer with 5 years experience building automated ML training and deployment pipelines using Kubeflow, MLflow, Triton Inference Server, PyTorch, Docker, and Kubernetes.

Technical Skills

MLOps & CI/CD: Kubeflow, MLflow, Triton Inference Server, Weights & Biases, DVC, Airflow
Cloud & Compute: Kubernetes, Docker, PyTorch, Python, C++, AWS EKS, CUDA

Work Experience

Senior MLOps Engineer - AI Pipeline Solutions (2021 - Present)
- Automated end-to-end model retraining and deployment pipeline on Kubernetes using Kubeflow and MLflow.
- Converted PyTorch models to ONNX and deployed on Triton Inference Server, reducing latency by 50%.

ML Infrastructure Engineer - DeepScale Inc (2019 - 2021)
- Built containerized model training infrastructure on AWS GPU clusters.

Education

M.S. in Computer Science - Stanford University